

```

In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

#Demonstrating Simple Linear Refression with 1 independent variable

df=pd.read_csv("data-Delhi.csv")

print(df)

X = np.array(df['Confirmed']).reshape(-1, 1)
y = np.array(df['Deaths']).reshape(-1, 1)

from sklearn.linear_model import LinearRegression
reg = LinearRegression().fit(X, y)
r2=reg.score(X, y)
coef=reg.coef_

intc=reg.intercept_

line=intc + coef*X

y_pred = reg.predict(y)

plt.scatter(X,y,color="r",marker="o",s=30)
plt.plot(X,line,color="y")
plt.xlabel('Cumulative Confirmed Cases')
plt.ylabel('Cumulative Confirmed Deaths')
plt.title('Covid19 February 2021 Delhi Data')
plt.show()

#To calculate standard error of regression slope
se = np.sqrt(1/(len(df)-2)*sum((y-y_pred)*(y-y_pred))/sum((X-np.mean(X))*(X-np.mean(X))))

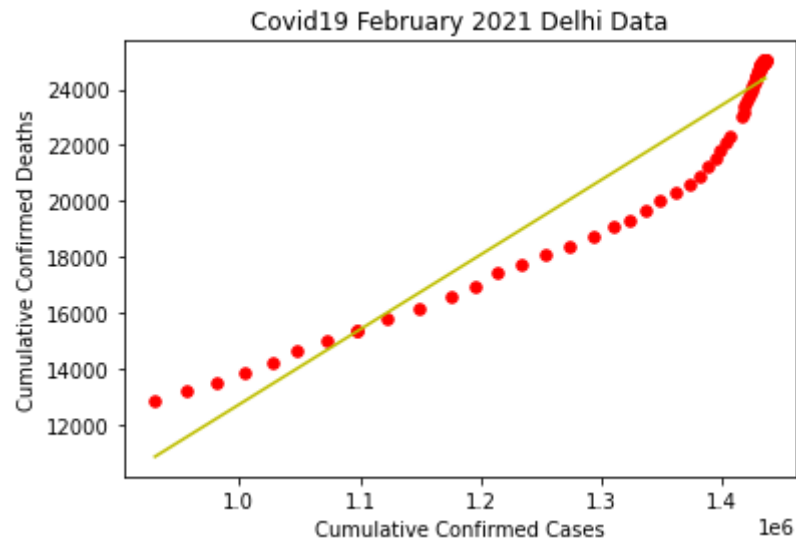
print("Intercept value = ",intc)
print("Slope value = ",coef[0])
print("r2 value = ",r2)
print("Standard Error of regression slope = ",se)

```

| | Unnamed: 0 | | Time | State | Confirmed | Recovered | Deaths | \ |
|----|------------|------------|----------|-------|-----------|-----------|--------|---|
| 0 | 0 | 2021-04-22 | 22:46:02 | Delhi | 930179 | 831928 | 12887 | |
| 1 | 1 | 2021-04-22 | 23:54:36 | Delhi | 956348 | 851537 | 13193 | |
| 2 | 2 | 2021-04-23 | 23:24:43 | Delhi | 980679 | 875109 | 13541 | |
| 3 | 3 | 2021-04-24 | 23:24:30 | Delhi | 1004782 | 897804 | 13898 | |
| 4 | 4 | 2021-04-25 | 22:29:17 | Delhi | 1027715 | 918875 | 14248 | |
| .. | ... | | ... | ... | ... | ... | ... | |
| 85 | 85 | 2021-07-11 | 19:02:20 | Delhi | 1435083 | 1409325 | 25015 | |
| 86 | 86 | 2021-07-12 | 19:58:16 | Delhi | 1435128 | 1409417 | 25018 | |
| 87 | 87 | 2021-07-14 | 13:56:28 | Delhi | 1435204 | 1409501 | 25020 | |
| 88 | 88 | 2021-07-14 | 18:39:07 | Delhi | 1435281 | 1409572 | 25021 | |
| 89 | 89 | 2021-07-15 | 19:34:29 | Delhi | 1435353 | 1409660 | 25022 | |

| | Active | New Cases |
|----|--------|-----------|
| 0 | 85364 | NaN |
| 1 | 91618 | 26169.0 |
| 2 | 92029 | 24331.0 |
| 3 | 93080 | 24103.0 |
| 4 | 94592 | 22933.0 |
| .. | ... | ... |
| 85 | 743 | 53.0 |
| 86 | 693 | 45.0 |
| 87 | 683 | 76.0 |
| 88 | 688 | 77.0 |
| 89 | 671 | 72.0 |

[90 rows x 8 columns]



Intercept value = $[-14037.2764932]$

Slope value = $[0.02677429]$

r2 value = 0.9313659523193988

Standard Error of regression slope = $[0.02892491]$

```
In [5]: import pandas as pd
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt

# set the max columns to none
pd.set_option('display.max_columns', None)

dp=pd.read_csv("insurance.csv")

df=dp
df.drop(df.index[(df['smoker']=='yes')],axis=0,inplace=True)
df.drop(df.index[(df['region']=='southeast')],axis=0,inplace=True)
df.drop(df.index[(df['region']=='southwest')],axis=0,inplace=True)
df.drop(df.index[(df['sex']=='female')],axis=0,inplace=True)
df.drop(df.index[(df['region']=='northwest')],axis=0,inplace=True)
df.drop('region',inplace=True,axis=1)
df.drop('sex',inplace=True,axis=1)
df.drop('smoker',inplace=True,axis=1)
y=df['charges']
X=df.drop(['charges'],axis=1)
```

```

print("\n")
print(X)
print("\n")

# importing train_test_split from sklearn
from sklearn.model_selection import train_test_split
# splitting the data
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 1)

# importing module
from sklearn.linear_model import LinearRegression
# creating an object of LinearRegression class
LR = LinearRegression()
# fitting the training data
LR.fit(x_train,y_train)

y_prediction = LR.predict(x_test)

print('Intercept = ')
print(LR.intercept_)

print('\nCoefficients = ')
print(LR.coef_)

print('\nr2 score for testing set = ', r2_score(y_test,y_prediction))

plt.scatter(y_test,y_prediction,color="g",marker="x",s=30)
plt.plot(y_test,y_test,color="r")
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.show()

plt.close()

plt.scatter(y_test,-y_test + LR.predict(x_test),color="blue",marker="o",s=10)
plt.plot(y_test,y_prediction-y_test)
plt.xlabel('Actual Price')
plt.ylabel('Residual Error')
plt.show()

print('\n')
print(y_prediction)

age=input("\n\nEnter age: ")

```

```

bmi=input("Enter bmi: ")
chil=input("Enter children: ")
print("\nInsurance Price as per users input:")
print(LR.predict(pd.DataFrame({'age':[age], 'bmi':[bmi], 'children':[chil]})))

```



| | age | bmi | children |
|------|-----|--------|----------|
| 8 | 37 | 29.830 | 2 |
| 10 | 25 | 26.220 | 0 |
| 17 | 23 | 23.845 | 0 |
| 44 | 38 | 37.050 | 1 |
| 60 | 43 | 27.360 | 3 |
| ... | ... | ... | ... |
| 1294 | 58 | 25.175 | 0 |
| 1296 | 18 | 26.125 | 0 |
| 1315 | 18 | 28.310 | 1 |
| 1318 | 35 | 39.710 | 4 |
| 1325 | 61 | 33.535 | 0 |

[125 rows x 3 columns]

Intercept =
-3809.9095124767628

Coefficients =
[279.14545162 21.57634111 822.22492955]

r2 score for testing set = 0.3110662695971691



```
[13736.17127083  2375.62672732 10910.25008812  8913.99459572
10122.33122643  2642.47366451 13034.39769918  2954.41515463
2059.20660796 10572.46928622 12719.86673361 12381.81818526
7997.10193108 11334.06480208  9475.03781125 14454.72198617
4337.30417526 14625.23056028  2647.75976324  3124.92372874
4770.72078037 10105.87751819 12385.70140246  4955.41702271
12032.44142879]
```

Enter age: 45
Enter bmi: 22.85
Enter children: 2

Insurance Price as per users input:
[10889.10506415]

```
In [6]: import pandas as pd
        from sklearn.metrics import r2_score
        import matplotlib.pyplot as plt

        # set the max columns to none
        pd.set_option('display.max_columns', None)

        dp=pd.read_csv("insurance.csv")

        dp=pd.concat([dp,pd.get_dummies(dp['sex'])],axis=1)
        dp=pd.concat([dp,pd.get_dummies(dp['smoker'])],axis=1)
        dp=pd.concat([dp,pd.get_dummies(dp['region'])],axis=1)

        dp.drop('region',inplace=True ,axis=1)
        dp.drop('sex',inplace=True ,axis=1)
        dp.drop('smoker',inplace=True ,axis=1)
        dp.drop('male',inplace=True ,axis=1)
        dp.drop('yes',inplace=True ,axis=1)
        dp.drop('southwest',inplace=True ,axis=1)

        print(dp)

        y=dp['charges']
        X=dp.drop(['charges'],axis=1)

        print("\n")
        print(X)
```

```

print("\n")

# importing train_test_split from sklearn
from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 1)

# importing module
from sklearn.linear_model import LinearRegression
# creating an object of LinearRegression class
LR = LinearRegression()
# fitting the training data
LR.fit(x_train,y_train)

print("\nIntercept = ")
print(LR.intercept_)

print("Coefficients = ")
print(LR.coef_)

y_prediction = LR.predict(x_test)

print('\nr2 score for testing set = ', r2_score(y_test,y_prediction))

plt.scatter(y_test,y_prediction,color="g",marker="x",s=30)
plt.plot(y_test,y_test,color="r")
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.show()

plt.close()

plt.scatter(y_test,-y_test + LR.predict(x_test),color="blue",marker="o",s=10)
plt.plot(y_test,y_prediction-y_prediction,)
plt.xlabel('Actual Price')
plt.ylabel('Residual Error')
plt.show()

print('\n')
print('Predicted values for testing set(y_prediction)\n')
print(y_prediction)

```


| | age | bmi | children | charges | female | no | northeast | northwest | \ |
|------|-----|--------|----------|-------------|--------|----|-----------|-----------|---|
| 0 | 19 | 27.900 | 0 | 16884.92400 | 1 | 0 | 0 | 0 | |
| 1 | 18 | 33.770 | 1 | 1725.55230 | 0 | 1 | 0 | 0 | |
| 2 | 28 | 33.000 | 3 | 4449.46200 | 0 | 1 | 0 | 0 | |
| 3 | 33 | 22.705 | 0 | 21984.47061 | 0 | 1 | 0 | 1 | |
| 4 | 32 | 28.880 | 0 | 3866.85520 | 0 | 1 | 0 | 1 | |
| ... | ... | ... | ... | ... | ... | .. | ... | ... | |
| 1333 | 50 | 30.970 | 3 | 10600.54830 | 0 | 1 | 0 | 1 | |
| 1334 | 18 | 31.920 | 0 | 2205.98080 | 1 | 1 | 1 | 0 | |
| 1335 | 18 | 36.850 | 0 | 1629.83350 | 1 | 1 | 0 | 0 | |
| 1336 | 21 | 25.800 | 0 | 2007.94500 | 1 | 1 | 0 | 0 | |
| 1337 | 61 | 29.070 | 0 | 29141.36030 | 1 | 0 | 0 | 1 | |

| | southeast |
|------|-----------|
| 0 | 0 |
| 1 | 1 |
| 2 | 1 |
| 3 | 0 |
| 4 | 0 |
| ... | ... |
| 1333 | 0 |
| 1334 | 0 |
| 1335 | 1 |
| 1336 | 0 |
| 1337 | 0 |



[1338 rows x 9 columns]

| | age | bmi | children | female | no | northeast | northwest | southeast |
|------|-----|--------|----------|--------|----|-----------|-----------|-----------|
| 0 | 19 | 27.900 | 0 | 1 | 0 | 0 | 0 | 0 |
| 1 | 18 | 33.770 | 1 | 0 | 1 | 0 | 0 | 1 |
| 2 | 28 | 33.000 | 3 | 0 | 1 | 0 | 0 | 1 |
| 3 | 33 | 22.705 | 0 | 0 | 1 | 0 | 1 | 0 |
| 4 | 32 | 28.880 | 0 | 0 | 1 | 0 | 1 | 0 |
| ... | ... | ... | ... | ... | .. | ... | ... | ... |
| 1333 | 50 | 30.970 | 3 | 0 | 1 | 0 | 1 | 0 |
| 1334 | 18 | 31.920 | 0 | 1 | 1 | 1 | 0 | 0 |
| 1335 | 18 | 36.850 | 0 | 1 | 1 | 0 | 0 | 1 |
| 1336 | 21 | 25.800 | 0 | 1 | 1 | 0 | 0 | 0 |
| 1337 | 61 | 29.070 | 0 | 1 | 0 | 0 | 1 | 0 |

[1338 rows x 8 columns]

```
Intercept =  
11343.689963450903  
Coefficients =  
[ 257.49024669 321.62189278 408.06102001 242.15306559  
 -23786.48604536 903.03300778 506.93644423 -135.34287187]
```

```
r2 score for testing set = 0.7623311844057112
```



Predicted values for testing set(y_prediction)

```
[ 4383.68089988 12885.03892192 12589.21653212 13286.22919217
   544.72832757 32117.58400779 12919.04237221 12318.62183013
  3784.29145555 29468.45725408 11002.8139431 17539.69473777
  8681.35471964 8349.04325528 3130.12725504 10445.83896118
 3863.74357865 6944.62510786 15009.63121084 14441.59911874
12543.65768867 32958.72553095 9072.63608136 8986.85860053
 3022.85773294 8164.97136102 9556.07558002 10743.20363927
 7694.01743692 4373.43771674 14140.93557984 5811.78545062
34631.91316718 27009.11191231 33348.14098668 9532.96786929
30421.65017927 26648.91186842 15157.78333287 33895.76121465
 6303.38552088 14059.15156303 10713.4467824 15089.36171493
 4187.95334069 13106.4297513 4336.19603407 28607.05556216
 7243.57117377 14269.4643165 13282.36924936 12329.61280721
 1851.87215658 8876.2837892 26089.18341811 10125.8221046
34218.77265378 14537.70022165 3232.07805794 5889.64309508
 6558.45711628 14952.73214832 26943.84457634 3272.57672674
15795.18877494 11220.12036023 11132.67761401 10461.51218201
 1520.17580687 25268.32319722 37555.4332681 33131.32070966
 1986.54437212 11348.45648105 13683.62487834 34970.76597049
 3194.05204265 3875.19388449 10355.84468565 10429.85383112
   -74.18168095 14069.96921025 10335.95235396 3160.49129709
33495.55139469 33108.38629603 7159.042252 37712.17792565
12860.01613403 10312.33535752 30118.39165257 33999.155218
14744.35977759 10797.48057723 228.32604517 10550.25751993
 9637.2654186 14963.62716464 14973.49438453 6077.52837971
13679.44499708 26048.6188477 28140.15460801 27428.44651929
35323.96326034 27120.17093173 635.73242244 9265.30720109
 4700.17995399 12458.33462103 5334.04136712 4797.80959774
 1053.28620015 18801.23368294 3268.21781045 1680.06692797
11731.45541277 12594.4560403 11876.24500234 3722.26917923
 8907.38977334 13909.79277731 7727.28039545 6573.92347482
36668.28291771 12172.54974158 12246.4759298 29298.69540744
36065.08836969 11635.06903459 28119.47917939 -420.5228157
 8255.48679122 31611.56891923 8278.51950655 -682.91733795
 1175.50251941 4610.52460783 7592.72365991 12602.74525758
14871.84794414 8696.2661006 28916.17140639 15712.12938325
14688.56307722 11117.34115616 1910.78149758 10065.51386262
 3785.83713249 6165.85822972 11400.42215978 5505.08475585
14580.76982237 13691.35579602 12694.51188244 7023.42319484
12388.68766385 10922.09183278 10269.55783904 4543.27270357
 5648.10144357 40390.9900769 13059.47316213 4308.66813543
```

| | | | |
|----------------|----------------|----------------|----------------|
| 8433.53823713 | 4680.92297563 | 32207.14761827 | 11261.09752853 |
| 10966.92628193 | 6893.83017801 | 6439.49932262 | 6698.81354717 |
| 33082.53354683 | 34892.66990169 | 2163.75212652 | 7664.10129233 |
| 5208.63123781 | 15537.4388228 | 1472.95942494 | 11431.38761905 |
| 13442.52462926 | 11497.84155642 | 10547.85065715 | 13216.06609157 |
| 2392.9275311 | 27535.86192673 | 2350.29363146 | 14750.02090702 |
| 6294.4943912 | 10590.51504221 | 14975.55458721 | 38857.75707767 |
| 2100.48817818 | 1489.62172706 | 5170.63120404 | 7556.77055613 |
| 7905.80683902 | 4503.61764622 | 10680.78553577 | 8938.12057203 |
| 9389.70713251 | 11104.75136012 | 10325.31689891 | 9247.40925093 |
| 8075.54835929 | 895.79174623 | 10136.82246673 | 7306.72664577 |
| 6626.07986045 | 11706.84936779 | 5409.99685749 | 32864.25315855 |
| 7088.39118065 | 6309.6941707 | 7934.10447803 | 38948.10610123 |
| 11941.19483711 | 28316.17975841 | 2882.4783976 | 33202.36401978 |
| 3690.60862539 | 31577.22772525 | 13825.53657174 | 2716.91852953 |
| 1908.80043495 | 1262.92212969 | 6109.40830379 | 4463.80387639 |
| 25580.05728181 | 15737.66640221 | 5345.8549026 | 13030.85900261 |
| 38954.05091304 | 4792.05740177 | 12711.42561622 | 11335.66208015 |
| 27785.54316341 | 2794.86874955 | 13392.79241645 | 5727.91540048 |
| 15215.43600554 | 5772.15783816 | 16929.82927411 | 3896.74375465 |
| 12197.3470759 | 34682.24329155 | 10666.53272796 | 10601.36016707 |
| 4875.20490336 | 16734.59399629 | 14399.64496923 | 5497.30018065 |
| 11149.82336777 | 12497.70437379 | 4626.74808217 | 7169.33486073 |
| 27667.13758601 | 32240.5545494 | -474.41779055 | 40306.05467371 |
| 9397.25562995 | 7750.27185181 | 10671.66257411 | 33555.1844395 |
| 35949.5230514 | 36650.46723087 | 4961.92884343 | 6116.92057448] |

In []: